

PAPER ANALYSIS



Presented by Yannis He

-

Paper: **VoxelNet: End-to-End Learning for Point Cloud Based 3D Object Detection**

Published Date: 17 Nov, 2017

Authors: Yin Zhou, Oncel Tuzel | Apple.Inc

State-of-the-art LiDAR based 3D detection methods on the KITTI car detection benchmark

<https://arxiv.org/abs/1711.06396>

- Background

- To interface a highly sparse LiDAR point cloud with a region proposal network (RPN), most existing efforts have focused on hand-crafted feature representations (Ex. Bird's eye view projection)
- Hand-crafted features yield satisfactory results when rich and detailed 3D shape info is available.
 - But unable to adapt to more complex shapes and scenes, and learn require invariances from data
- Image-based (2D input) 3D detection approaches are bounded by the accuracy of depth estimation

- Contribution:

- Remove the need of manual feature engineering for 3D point clouds: VoxelNet, a generic 3D detection network that unifies feature extraction and bounding box prediction into a single stage, end-to-end trainable deep network.
 - Divides a point cloud into equally spaced 3D voxels
 - Transforms a group of points within each voxel into a unified feature representation through the newly introduced voxel feature encoding (VFE) layer.
 - I.e. point cloud is encoded as descriptive volumetric representation, which is then connected to a RPN to generate detections
- State-of-the-art LiDAR based 3D detection methods on the KITTI car detection benchmark

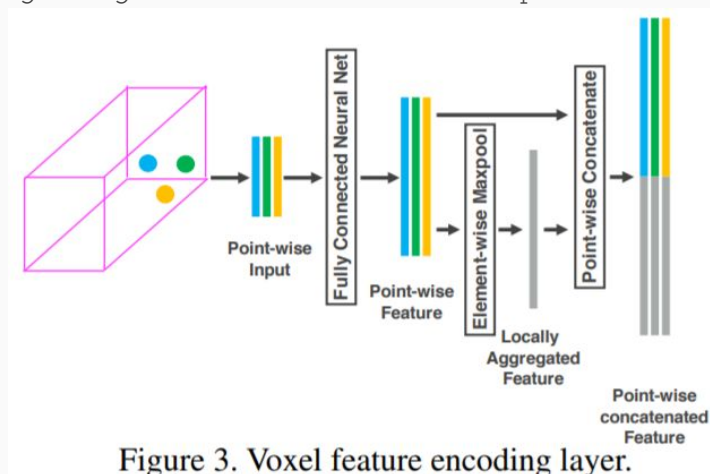


Figure 3. Voxel feature encoding layer.

- Background:
 - Comparing with images, LiDAR provides reliable depth information that can be used to accurately localize objects and characterize their shape
 - But LiDAR points are sparse and have highly variable point density due to factors as non-uniform sampling of the 3D space, effective range of sensors, occlusion, relative pose.
 - This results that many approaches manually crafted feature representations for point clouds that are tuned for 3D object detection.
 - Ex. project point clouds into a perspective view and apply image-based feature extraction.
 - However, manual design choices introduce an information bottleneck, which prevents effectively exploiting 3D shape information and the required invariance for detection task.
- Contribution:
 - 3D detection framework that simultaneously learns a discriminative feature representation from point clouds and predicts accurate 3D bounding boxes, in an end-to-end fashion.
 - Voxel feature encoding (VFE) layer, which enables inter-point interaction within a voxel by combining point-wise features with a locally aggregated feature
 - Stacking VFE layers allows learning complex features for characterizing local 3D shape information
 - VoxelNet divides point cloud into equally spaced 3D voxels, encodes each voxel via stacked VFE layers, and then 3D convolution further aggregates local voxel features, transforming the point cloud into a high-dimensional volumetric representation. Finally a RPN consumes the volumetric representation and yields the detection result.
 - Benefits both from sparse point structure and efficient parallel processing on voxel grid

- 3 functional blocks:
 - Feature learning network
 - Convolutional middle layers
 - Region proposal network

ARCHITECTURE - VoxelNet

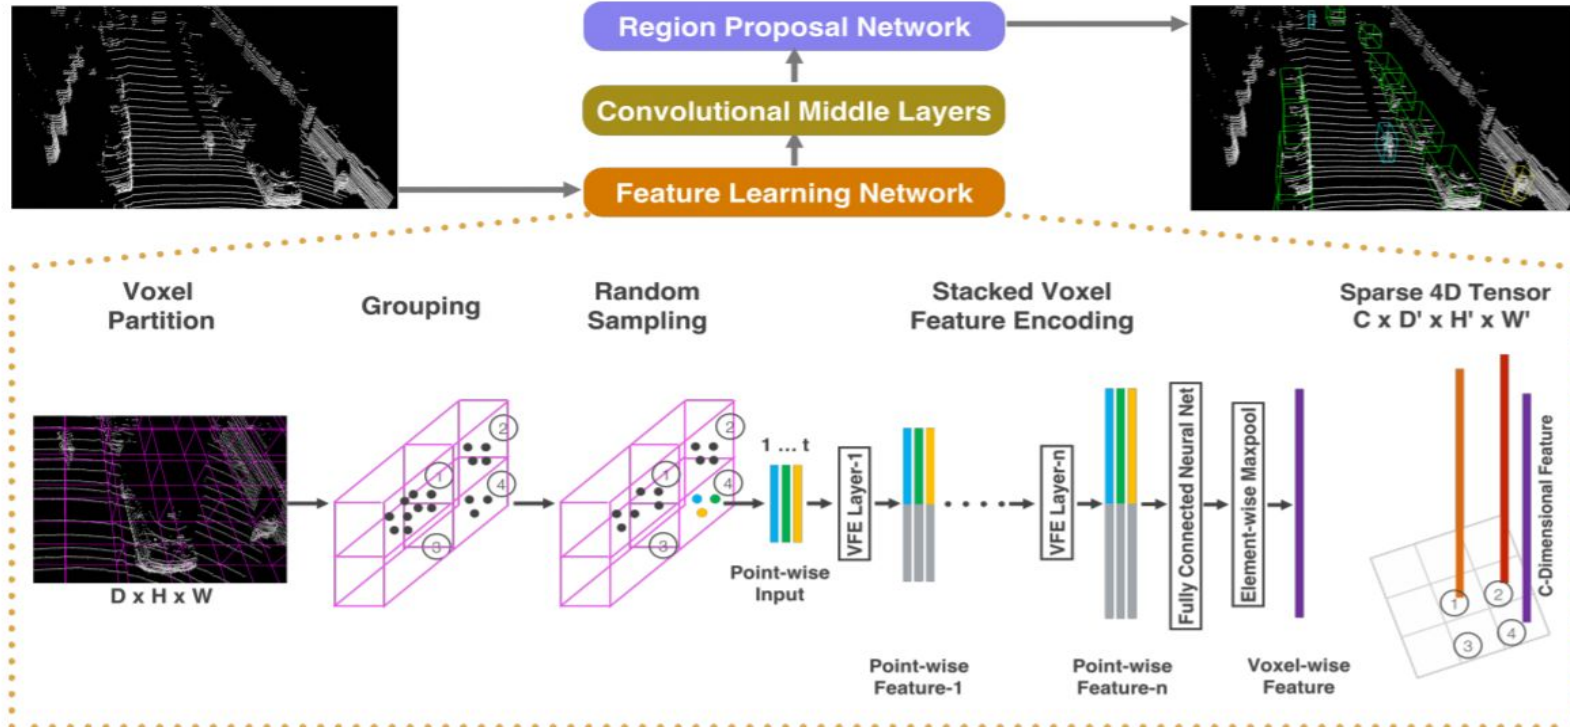



Figure 2. VoxelNet architecture. The feature learning network takes a raw point cloud as input, partitions the space into voxels, and transforms points within each voxel to a vector representation characterizing the shape information. The space is represented as a sparse 4D tensor. The convolutional middle layers processes the 4D tensor to aggregate spatial context. Finally, a RPN generates the 3D detection.


ARCHITECTURE - VoxelNet: FEATURE LEARNING NETWORK

- 1st block: Feature learning network:
 - **Voxel Partition:**
 - Given a point cloud, subdivide 3D space into equalled spaced voxels.
 - **Grouping:**
 - Group points according to the voxel they reside in.
 - Since LiDAR point cloud is sparse and has highly variable point density, voxels will contain a variable number of points after grouping.
 - **Random Sampling:**
 - Randomly sample a fixed number, T , of points from voxels containing more than T points
 - 1) computational saving &
 - 2) decreases the imbalance (sampling bias) of points between voxels and add variation




D x H x W

Voxel Partition



Grouping



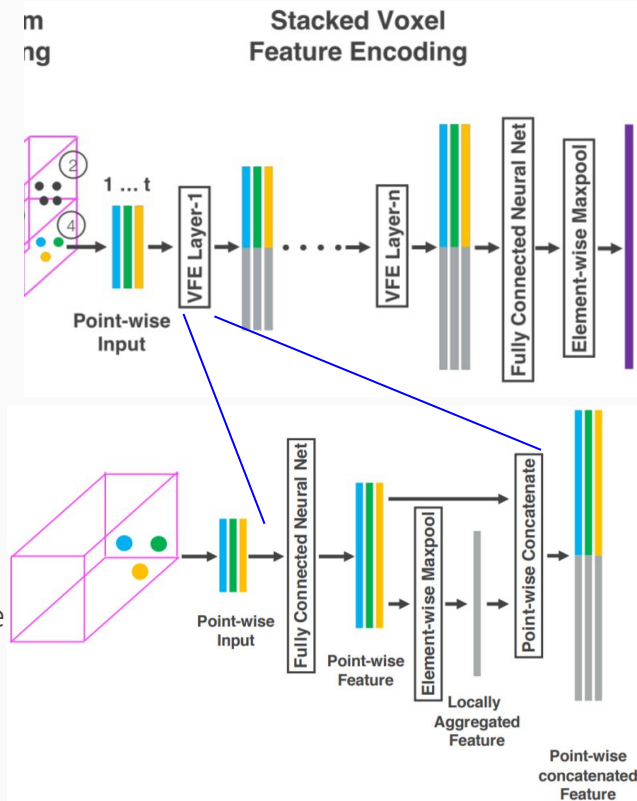
Random Sampling

ARCHITECTURE - VoxelNet: FEATURE LEARNING NETWORK (CONT')

- 1st block: Feature learning network (Cont'):


- Stacked Voxel Feature Encoding:**

- The key innovation: the chain of VFE layers
 - 1. Compute local mean as the **centroid** (v_x, v_y, v_z) of all points in a voxel, $V = \{p_i = [x_i, y_i, z_i, r_i]^T \in \mathbb{R}^4\}_{i=1, \dots, t}$
 - 2. When augment each point with the relative offset w.r.t. to the centroid and **obtain input features** $V_{in} = \{\hat{p}_i = [x_i, y_i, z_i, r_i, x_i - v_x, y_i - v_y, z_i - v_z]^T \in \mathbb{R}^7\}_{i=1, \dots, t}$.
 - 3. Each \hat{p}_i is transformed through the fully connected network (FCN) into a feature space, where we can aggregate information from the point features $f_i \in \mathbb{R}^m$ to encode the shape of the surface contained within the voxel. \rightarrow **point-wise feature representation**
 - a. FCN: linear layer, batch normalization (BN), ReLU
 - 4. **Element-wise MaxPooling** across all f_i associate to V to get the locally **aggregated feature** $\tilde{f} \in \mathbb{R}^m$ for V
 - 5. Augment each f_i with \tilde{f} to form the point-wise concatenated feature as $f_i^{out} = [f_i^T, \tilde{f}^T]^T \in \mathbb{R}^{2m} \rightarrow$ output feature: $V_{out} = \{f_i^{out}\}_{i=1, \dots, t}$
 - Output feature combines both point-wise features and locally aggregated feature, stacking VFE layers encodes point interactions with a voxel and enables the final feature representation to learn descriptive shape information.



ARCHITECTURE - VOXELNET: FEATURE LEARNING NETWORK (CONT')

- 1st block: Feature learning network (Cont'):
 - **Sparse Tensor Representation::**
 - By processing only non-empty voxels, we obtain a list of voxel features
 - Each uniquely associated to the spatial coordinates of a particular non-empty voxel.
 - 90% of voxels typically are empty, out of ~100k points
 - The obtained list of voxel-wise features can be represented as a sparse 4D tensor, of size $C \times D' \times H' \times W'$ as shown in the figure
 - Representing non-empty voxel features as a sparse tensor greatly reduces the memory usage and computation cost during backpropagation..




*Sparse 4D tensor
 $C \times D' \times H' \times W'$*

ARCHITECTURE - VOXELNET: CONVOLUTIONAL MIDDLE LAYERS

- 2nd block: Convolutional Middle Layers
 - $\text{ConvMD}(c_{\text{in}}, c_{\text{out}}, k, s, p)$ is used to represent an M-dimensional convolution operator, where c_{in} and c_{out} are the number of input and output channels. k, s, p are the M-dimensional vectors corresponding to kernel size, stride size, and padding size respectively.
 - Each convolutional middle layer applies 3D convolution, BN layer, and ReLU layer sequentially.
 - Convolutional middle layers aggregate voxel-wise features within a progressively expanding receptive field, adding more context to the shape description.

ARCHITECTURE - VOXELNET: REGION PROPOSAL NETWORK

- 3rd block: Region Proposal Network
 - An important building block of top-performing object detection framework
 - In this work, RPN has several key modification and combined with the feature learning network and convolutional middle layers to form an end-to-end trainable pipeline
 - 3 blocks of fully convolutional layers:
 - Feature maps is mapped to the desired learning targets: 1) a probability score map & 2) a regression map



● Notation:

- Let $\{a_i^{\text{pos}}\}_{i=1\dots N_{\text{pos}}}$ be the set of N_{pos} positive anchors & $\{a_i^{\text{neg}}\}_{i=1\dots N_{\text{neg}}}$ be the set of N_{neg} positive anchors

- 3D ground truth box, $(x_c^g, y_c^g, x_c^g, l^g, w^g, h^g, \theta^g)$

- x_c^g, y_c^g, z_c^g represent center location
- l^g, w^g, h^g are length, width, height of the box
- θ^g is the yaw rotation around Z-axis

- Positive anchor, $(x_c^a, y_c^a, x_c^a, l^a, w^a, h^a, \theta^a)$

- Residual vector, $\mathbf{u}^* \in \mathbb{R}^7$, containing 7 regression targets corresponding to

- center location $\Delta x, \Delta y, \Delta z$
- 3 dimension: $\Delta l, \Delta w, \Delta h$
- Rotation $\Delta \theta$

- Diagonal of the base of the anchor box: $d^a = \sqrt{(l^a)^2 + (w^a)^2}$

- Loss function:

$$L = \alpha \frac{1}{N_{\text{pos}}} \sum_i L_{\text{cls}}(p_i^{\text{pos}}, 1) + \beta \frac{1}{N_{\text{neg}}} \sum_j L_{\text{cls}}(p_j^{\text{neg}}, 0) \quad \text{where } \alpha, \beta \text{ are positive constants}$$

$$+ \frac{1}{N_{\text{pos}}} \sum_i L_{\text{reg}}(\mathbf{u}_i, \mathbf{u}_i^*) \quad (2)$$

- p_i^{pos} & p_j^{neg} represent the softmax output for positive anchor a_i^{pos} and negative anchor a_j^{neg} respectively
- $\mathbf{u}_i \in \mathbb{R}^7$ and $\mathbf{u}^* \in \mathbb{R}^7$ are the regression output and ground truth for positive anchor a_i^{pos} .
- L_{cls} : binary cross entropy loss & L_{reg} is the regression loss (where SmoothL1 function is used)

| Method | Modality | Car | | | Pedestrian | | | Cyclist | | |
|--------------------|------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| | | Easy | Moderate | Hard | Easy | Moderate | Hard | Easy | Moderate | Hard |
| Mono3D [3] | Mono | 5.22 | 5.19 | 4.13 | N/A | N/A | N/A | N/A | N/A | N/A |
| 3DOP [4] | Stereo | 12.63 | 9.49 | 7.59 | N/A | N/A | N/A | N/A | N/A | N/A |
| VeloFCN [22] | LiDAR | 40.14 | 32.08 | 30.47 | N/A | N/A | N/A | N/A | N/A | N/A |
| MV (BV+FV) [5] | LiDAR | 86.18 | 77.32 | 76.33 | N/A | N/A | N/A | N/A | N/A | N/A |
| MV (BV+FV+RGB) [5] | LiDAR+Mono | 86.55 | 78.10 | 76.67 | N/A | N/A | N/A | N/A | N/A | N/A |
| HC-baseline | LiDAR | 88.26 | 78.42 | 77.66 | 58.96 | 53.79 | 51.47 | 63.63 | 42.75 | 41.06 |
| VoxelNet | LiDAR | 89.60 | 84.81 | 78.57 | 65.95 | 61.05 | 56.98 | 74.41 | 52.18 | 50.49 |

Table 1. Performance comparison in bird's eye view detection: average precision (in %) on KITTI validation set.

| Method | Modality | Car | | | Pedestrian | | | Cyclist | | |
|--------------------|------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| | | Easy | Moderate | Hard | Easy | Moderate | Hard | Easy | Moderate | Hard |
| Mono3D [3] | Mono | 2.53 | 2.31 | 2.31 | N/A | N/A | N/A | N/A | N/A | N/A |
| 3DOP [4] | Stereo | 6.55 | 5.07 | 4.10 | N/A | N/A | N/A | N/A | N/A | N/A |
| VeloFCN [22] | LiDAR | 15.20 | 13.66 | 15.98 | N/A | N/A | N/A | N/A | N/A | N/A |
| MV (BV+FV) [5] | LiDAR | 71.19 | 56.60 | 55.30 | N/A | N/A | N/A | N/A | N/A | N/A |
| MV (BV+FV+RGB) [5] | LiDAR+Mono | 71.29 | 62.68 | 56.56 | N/A | N/A | N/A | N/A | N/A | N/A |
| HC-baseline | LiDAR | 71.73 | 59.75 | 55.69 | 43.95 | 40.18 | 37.48 | 55.35 | 36.07 | 34.15 |
| VoxelNet | LiDAR | 81.97 | 65.46 | 62.85 | 57.86 | 53.42 | 48.87 | 67.17 | 47.65 | 45.11 |

Table 2. Performance comparison in 3D detection: average precision (in %) on KITTI validation set.

RESULT (CONT')

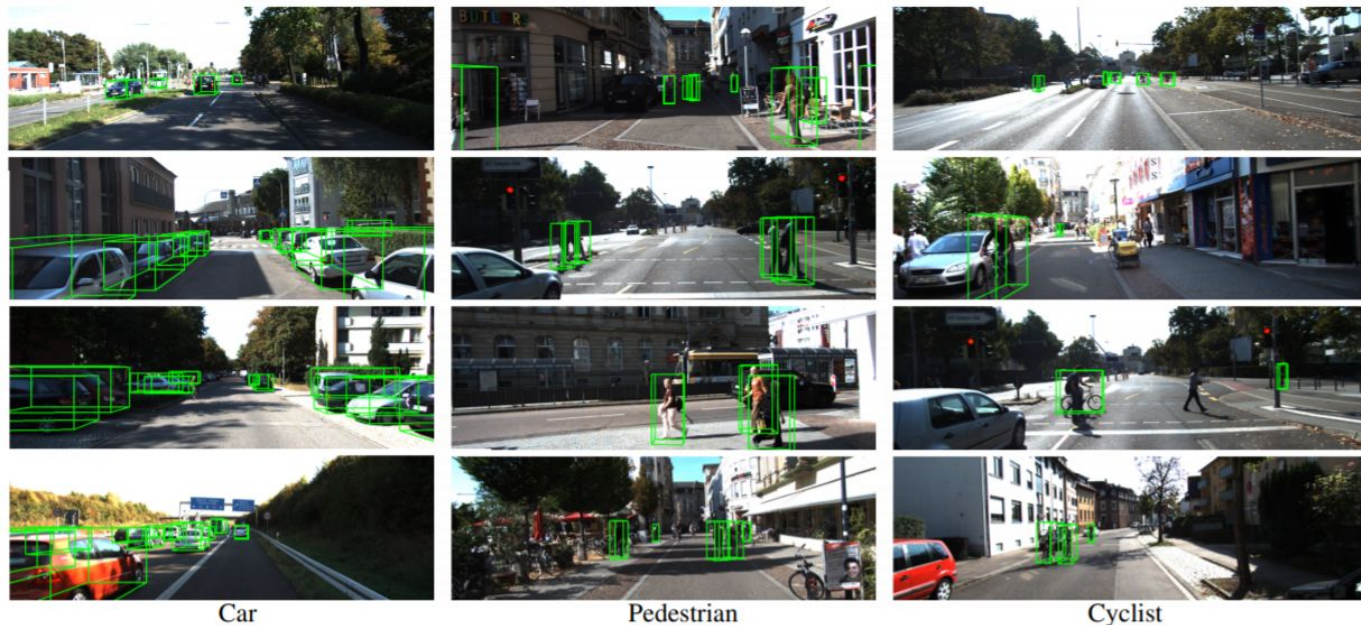


Figure 6. Qualitative results. For better visualization 3D boxes detected using LiDAR are projected on to the RGB images.

| Benchmark | Easy | Moderate | Hard |
|------------------------------|-------|----------|-------|
| Car (3D Detection) | 77.47 | 65.11 | 57.73 |
| Car (Bird's Eye View) | 89.35 | 79.26 | 77.39 |
| Pedestrian (3D Detection) | 39.48 | 33.69 | 31.51 |
| Pedestrian (Bird's Eye View) | 46.13 | 40.74 | 38.11 |
| Cyclist (3D Detection) | 61.22 | 48.36 | 44.37 |
| Cyclist (Bird's Eye View) | 66.70 | 54.76 | 50.55 |

Table 3. Performance evaluation on KITTI test set.